

# Linear Regression and Stochastic Gradient Descent

Winglok

2018 年 11 月 1 日

Download the dataset “housing\_scale”:

```
In [1]: import requests
        r = requests.get(''https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/regression/housing_scale'')
```

Load the dataset to X,y:

```
In [2]: from sklearn.datasets import load_svmlight_file
        from io import BytesIO
        X, y = load_svmlight_file(f=BytesIO(r.content), n_features=13)
        X = X.toarray()
```

Preprocess, change the shape of X and y:

```
In [3]: import numpy
        n_samples, n_features = X.shape
        X = numpy.column_stack((X, numpy.ones((n_samples, 1))))
        y = y.reshape((-1, 1))
```

Devide the dataset into traning set and validation set(traning size is 3/4,and validation size is rest here):

```
In [4]: from sklearn.model_selection import train_test_split
        X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.25)
```

Define regular term coefficients learning rate and max epoches:

```
In [5]: penalty_factor = 0.01
        learning_rate = 0.001
        max_epoch = 3500
        losses_train = []
        losses_val = []
```

Initialize w by different ways(using all zeros here):

```
In [6]: # select different initializing method
        w = numpy.zeros((n_features + 1, 1)) # initialize with zeros
        # w = numpy.random.random((n_features + 1, 1)) # initialize with random numbers
        # w = numpy.random.normal(1, 1, size=(n_features + 1, 1)) # initialize with zero normal distribution
```

Closed solution to linear regression:

$$\hat{w} = (X^T X)^{-1} X^T y$$

```
In [7]: X_train_t=X_train.transpose()
        temp1=numpy.dot(X_train_t,X_train)
        temp2=numpy.dot(numpy.linalg.inv(temp1),X_train_t)
        w=numpy.dot(temp2,y_train)
```

Loss function:

$$L = \frac{1}{2m} \sum_{i=1}^m (y_i - X_i \omega)^2$$

or

$$L = \frac{1}{2m} ||y_i - X_i \omega||_2^2$$

```
In [8]: Y_predict_closed = numpy.dot(X_train, w) # predict under the train set
        loss_train_closed = 1/2*numpy.average(numpy.square(Y_predict_closed-y_train))

        Y_predict_closed = numpy.dot(X_val, w) # predict under the validation set
        loss_val_closed = 1/2*numpy.average(numpy.square(Y_predict_closed-y_val))

        Y_predict_closed = numpy.dot(X, w) # under all dataset
        loss_closed=1/2*numpy.average(numpy.square(Y_predict_closed-y))
        print("The loss under the train set is ",loss_train_closed)
        print("The loss under the validation set is ",loss_val_closed)
        print("The loss under all dataset is ",loss_closed)
```

The loss under the train set is 11.366724927728239  
The loss under the validation set is 10.310000043399386  
The loss under all dataset is 11.101499512096295

### Train and predict by Stochastic Gradient Descent(SGD):

Here are fomulas used below:

we have loss function with L1 regularization (Lasso regression) :

$$L = \frac{1}{2m} \sum_{i=1}^m (y_i - X_i \omega)^2 + \frac{\lambda}{m} \|\omega\|_1$$

or

$$L = \frac{1}{2m} \|y_i - X_i \omega\|_2^2 + \frac{\lambda}{m} \|\omega\|_1$$

the gradient is:

$$G = \frac{\partial L}{\partial \omega} = -\frac{1}{m} \left( \sum_{i=1}^m y_i X_i^T - \left( \sum_{i=1}^m X_i^T X_i \right) \omega - \lambda \operatorname{sgn}(\omega) \right)$$

Details about the above derivation:

for above formula, X is a  $m \times 13$  matrix, and  $X_i$  is a row vector which represents a sample with 13 properties, then  $\omega$  is a  $13 \times 1$  matrix which from  $\omega_1$  to  $\omega_{13}$ , so we can seek partial guidance on  $\omega_1$  (Without loss of generality, We regard that  $X: m \times n$   $\omega: n \times 1$   $y: m \times 1$ ):

$$\begin{aligned} \frac{\partial L}{\partial \omega_1} &= \frac{\partial \frac{1}{2m} \sum_{i=1}^m (y_i - X_i \omega)^2 + \frac{\lambda}{m} \|\omega\|_1}{\partial \omega_1} \\ &= \frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) \frac{\partial \sum_{i=1}^m (y_i - X_i \omega)}{\partial \omega_1} + \frac{\lambda}{m} \operatorname{sgn}(\omega_1) \\ &= \frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) \frac{\partial (y_1 - X_{11}\omega_1 - X_{12}\omega_2 - \dots - X_{1n}\omega_n + y_2 - X_{21}\omega_1 - X_{22}\omega_2 - \dots - X_{2n}\omega_n + \dots + y_m - X_{m1}\omega_1 - X_{m2}\omega_2 - \dots - X_{mn}\omega_n)}{\partial \omega_1} + \frac{\lambda}{m} \operatorname{sgn}(\omega_1) \\ &= \frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) (-X_{i1} - X_{i2} - \dots - X_{in}) + \frac{\lambda}{m} \operatorname{sgn}(\omega_1) \\ &= -\frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) \sum_{j=1}^n X_{ij} + \frac{\lambda}{m} \operatorname{sgn}(\omega_1) \end{aligned}$$

We can also get follow formula and so on:

$$\frac{\partial L}{\partial \omega_2} = -\frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) \sum_{j=1}^n X_{ij} + \frac{\lambda}{m} \operatorname{sgn}(\omega_2)$$

thus:

$$\frac{\partial L}{\partial \omega} = \begin{bmatrix} \frac{\partial L}{\partial \omega_1} \\ \frac{\partial L}{\partial \omega_2} \\ \dots \\ \frac{\partial L}{\partial \omega_n} \end{bmatrix}$$

and

$$\begin{bmatrix} \sum_{j=1}^m X_{j1} \\ \sum_{j=1}^m X_{j2} \\ \dots \\ \sum_{j=1}^m X_{jn} \end{bmatrix} = \sum_{i=1}^m X_i^T$$

and finally we get:

$$\begin{aligned} \frac{\partial L}{\partial \omega} &= -\frac{1}{m} \sum_{i=1}^m (y_i - X_i \omega) X_i^T + \frac{\lambda}{m} \operatorname{sgn}(\omega) \\ &= -\frac{1}{m} \left( \sum_{i=1}^m y_i X_i^T - \left( \sum_{i=1}^m X_i^T X_i \right) \omega - \lambda \operatorname{sgn}(\omega) \right) \end{aligned}$$

update

$$\omega = \omega - \eta G$$

pre process:

```
In [9]: def sgn(v):
        num=0
        for i in v:
            if i>0:
                v[num]=1
            if i<0:
                v[num]=-1
            if i==0:
                v[num]=0
            num+=1
```

training and iteration:

```
In [10]: import random
w = numpy.zeros((n_features + 1, 1)) # initialize with zeros
for epoch in range(max_epoch):
    r=random.randint(0,X_train.shape[0]-1)
    X_train_s=X_train[r]
    y_train_s=y_train[r]
    X_train_s_t=X_train_s.reshape(X_train_s.shape[0],1)#choose one sample

    G = -((y_train_s-numpy.dot(X_train_s,w))*(X_train_s_t)-penalty_factor*w)# calculate the gradient
    G = -G
    w += learning_rate * G # update the parameters

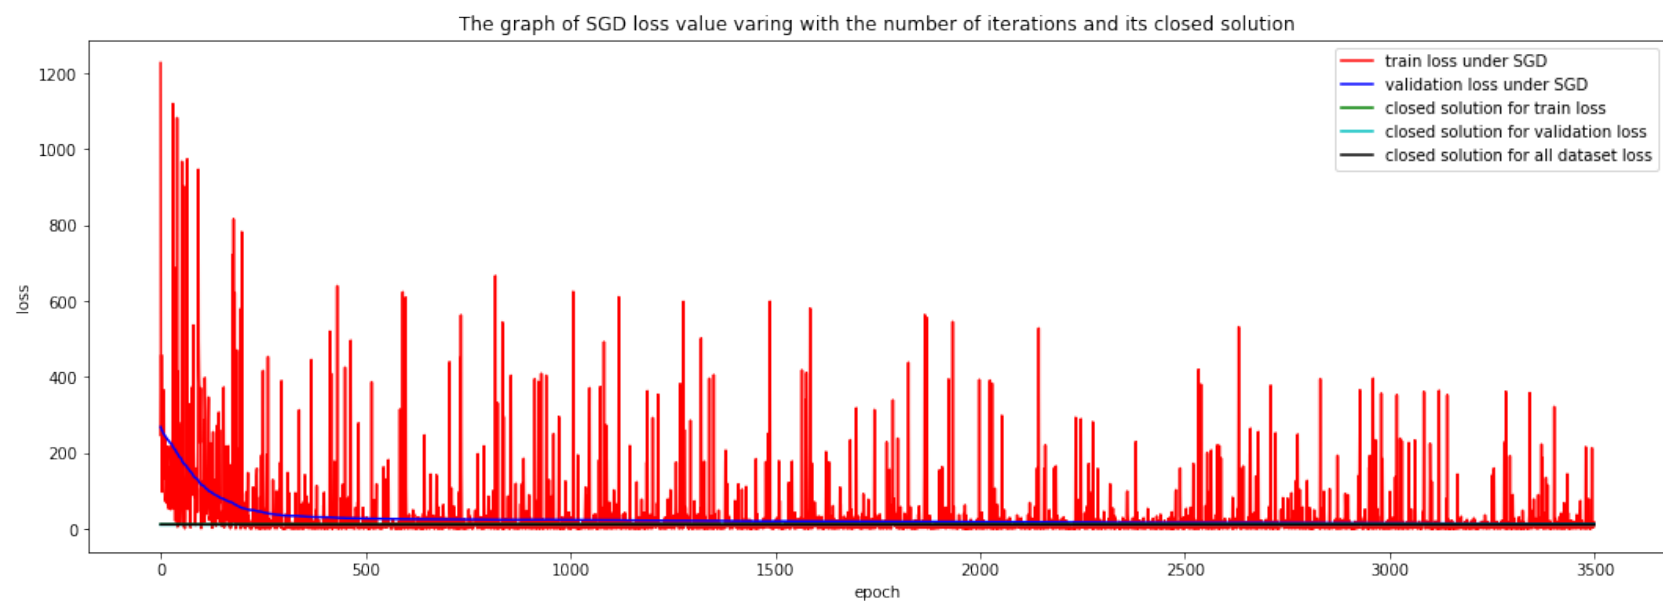
    Y_predict = numpy.dot(X_train_s, w) # predict under the train set
    loss_train = 1/2*(numpy.square(Y_predict-y_train_s))# calculate
    losses_train.append(loss_train)

    Y_predict = numpy.dot(X_val, w) # predict under the validation set
    loss_val = 1/2*numpy.average(numpy.square(Y_predict-y_val)) # calculate
    losses_val.append(loss_val)
```

Plot loss values with the number of iterations:

```
In [11]: %matplotlib inline
import matplotlib.pyplot as plt
losses_train_closed=[]
losses_val_closed=[]
losses_closed=[]
plt.figure(figsize=(18, 6))
plt.plot(losses_train, "-", color="r", label="train loss under SGD")
plt.plot(losses_val, "-", color="b", label="validation loss under SGD")
for i in range(3500):
    losses_train_closed.append(loss_train_closed)
    losses_val_closed.append(loss_val_closed)
    losses_closed.append(loss_closed)
plt.plot(losses_train_closed, "-", color="g", label="closed solution for train loss")
plt.plot(losses_val_closed, "-", color="c", label="closed solution for validation loss")
plt.plot(losses_closed, "-", color="k", label="closed solution for all dataset loss")
plt.xlabel("epoch")
plt.ylabel("loss")
plt.legend()
plt.title("The graph of SGD loss value varing with the number of iterations and its closed solution")
```

```
Out[11]: Text(0.5,1,'The graph of SGD loss value varing with the number of iterations and its closed solution')
```



#### References:

1. 机器学习之正则化（Regularization）[EB/OL]. <https://www.cnblogs.com/jianxinzhou/p/4083921.html>.
2. 回归系列之 L1 和 L2 正则化 [EB/OL]. <https://www.jianshu.com/p/a47c46153326>.
3. 机器学习总结 (一): 线性回归、岭回归、Lasso 回归 [EB/OL]. <https://blog.csdn.net/hzw19920329/article/details/77200475>.
4. 正则化为什么能防止过拟合 [EB/OL]. <https://www.cnblogs.com/alexanderkun/p/6922428.html>.
5. 最优化方法: 梯度下降（批梯度下降和随机梯度下降）[EB/OL]. <https://blog.csdn.net/pipisorry/article/details/23692455>.
6. 梯度下降、随机梯度下降和批量梯度下降 [EB/OL]. <https://www.cnblogs.com/louyihang-loves-baiyan/p/5136447.html>.
7. Markdown 数学公式 [EB/OL]. <http://blog.lisp4fun.com/2017/11/01/formula>.
8. Markdown 公式编辑学习笔记 [EB/OL]. <https://www.cnblogs.com/q735613050/p/7253073.html>.
9. NumPy 中文文档 [EB/OL]. <https://www.numpy.org.cn/>.